Building a “Hits” Predictive Model

Shelmith Nyagathiri Kariuki

Data Description

- The challenge was to build a predictive model that predicts the number of hits per session.

- The data consisted of 988,681 records and 10 variables.

Variable	Description
row_num	A number uniquely identifying each row.
locale	The platform of the session.
day_of_week	Mon-Fri, the day of the week of the session.
hour_of_day	00-23, the hour of the day of the session.
agent_id	The device used for the session.
entry_page	Describes the landing page of the session.
path_id_set	Shows all the locations that were visited during the session.
traffic_type	Indicates the channel the user cane through eg. search engine, email, ...
session_duration	The duration in seconds of the session.
hits	The number of interactions with the trivago page during the session.

Analysis Process

Cleaning

Converting string variables to factors, and numeric variables to actual numerics

Replacing “\N” with actual NAs

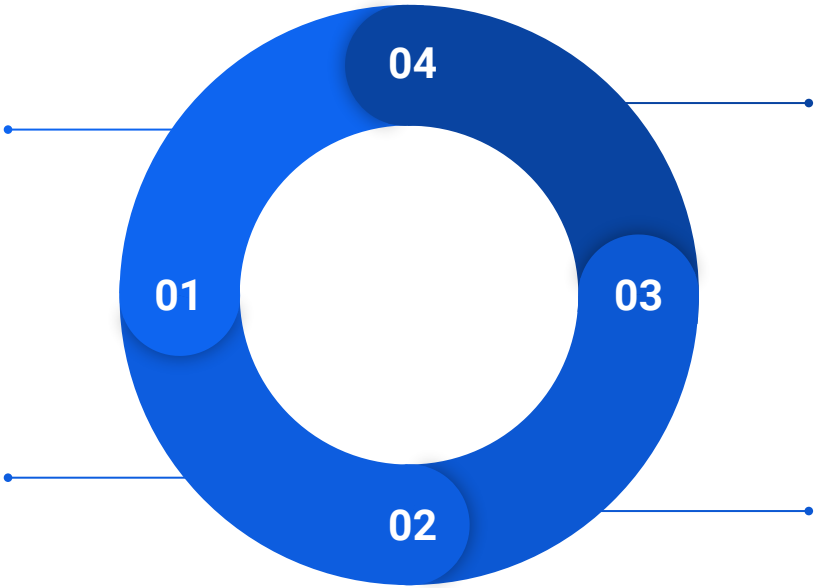

Modeling

Fitting the SVM and Random Forest models, and obtaining predicted values.

Feature Engineering

Generating additional variables from the ones that we already have

Data Exploration

Assessing relationships among the variables, and between the variables and the dependent variable

Feature Engineering

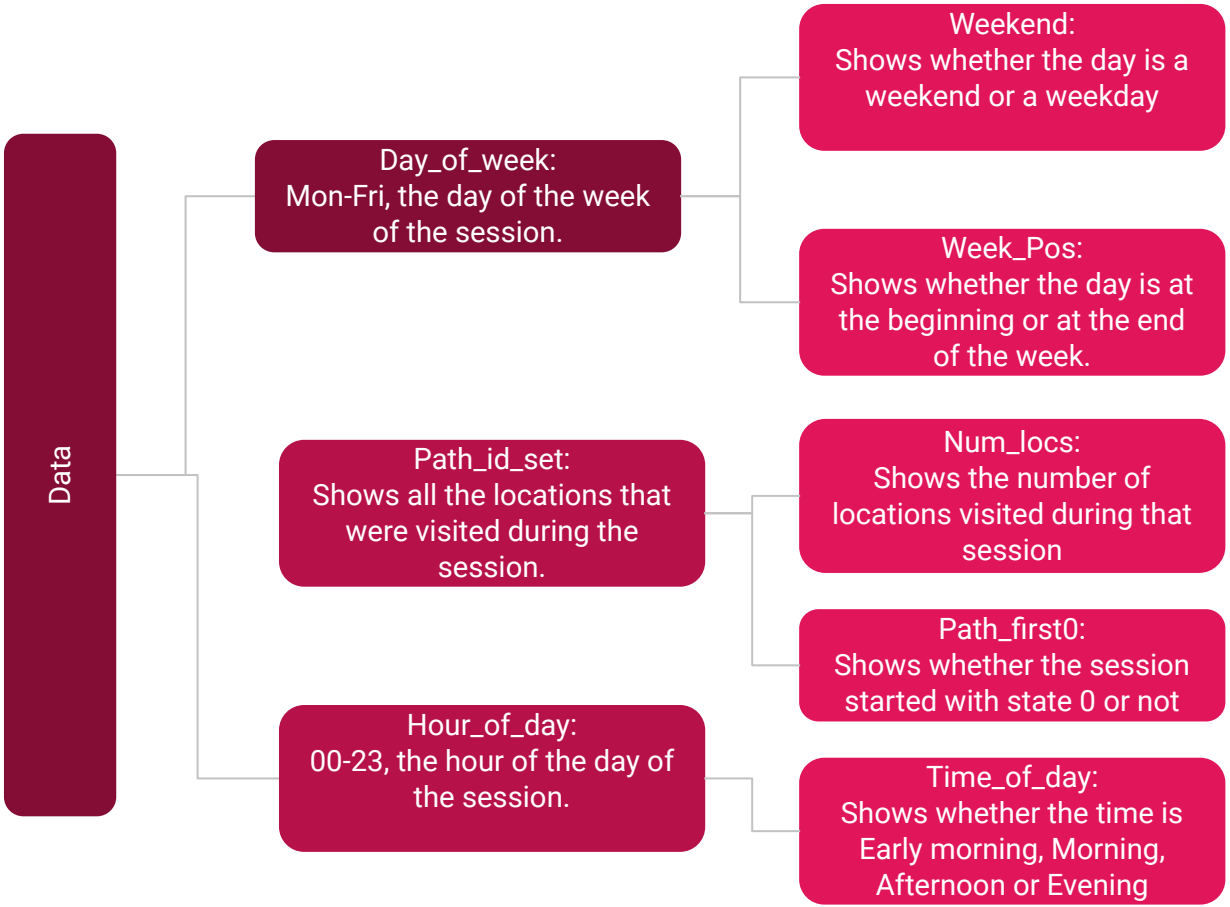

Data Exploration

Number of Hits

- A majority of visitors on the Trivago website make an average of about 5-10 hits per session, as seen on the third graph.

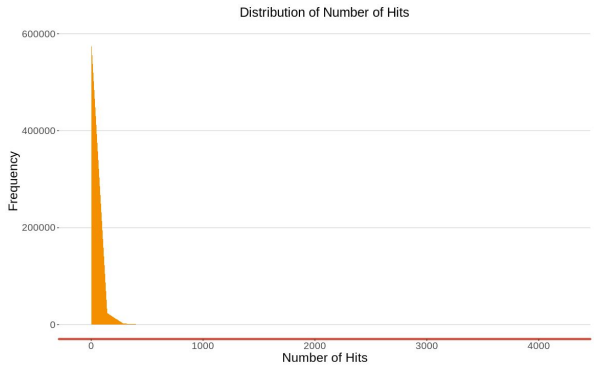

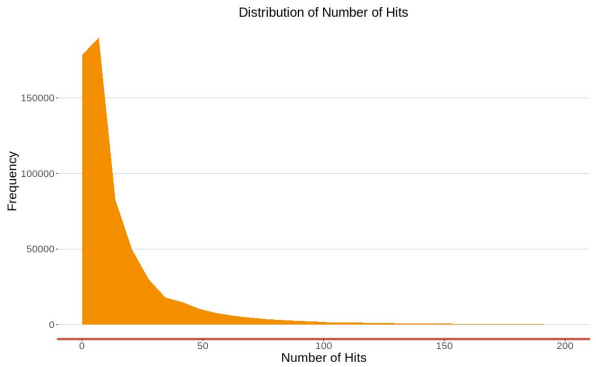

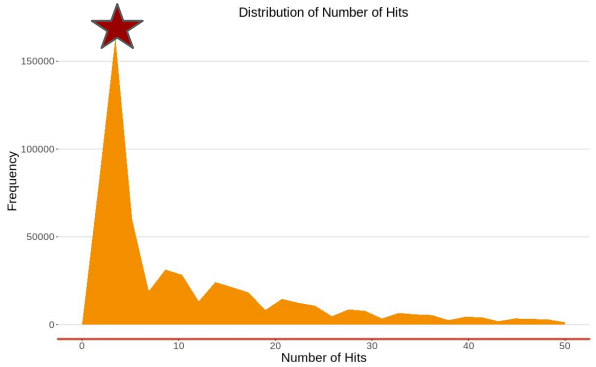

The whole dataset

Data truncated to only 200 hits

Data truncated to only 50 hits

Days of the week

- There is an upward trend in the average number of hits as the week progresses, but this decreases as the week comes to an end.

- There is a **significantly** higher number of hits during the weekdays, as opposed to during the weekends.

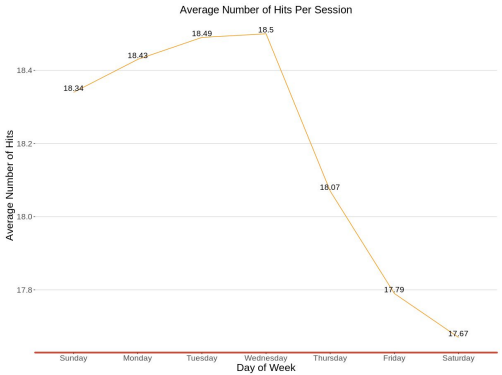

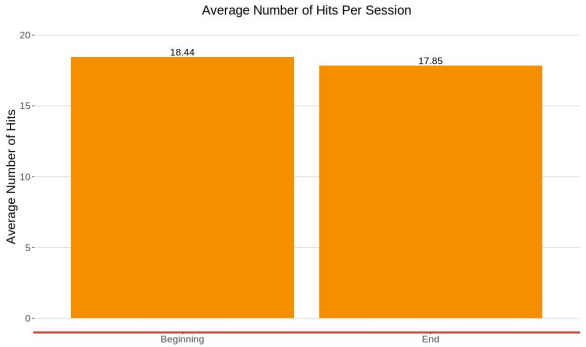

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
weekend	1	7016	7016	4.589	0.0322 *
Residuals	619233	946669561	1529		

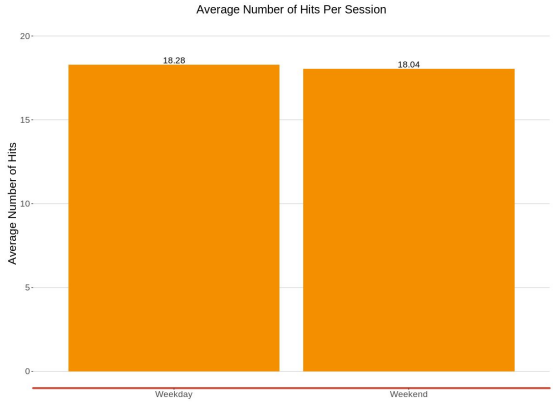

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
week_pos	1	51318	51318	33.57	0.00000000688 ***
Residuals	619233	946625260	1529		

Time of day

- The number of hits is relatively high at the beginning of each day.
- There is a very low negative correlation (0.00796667) between the number of hits and the time of day.

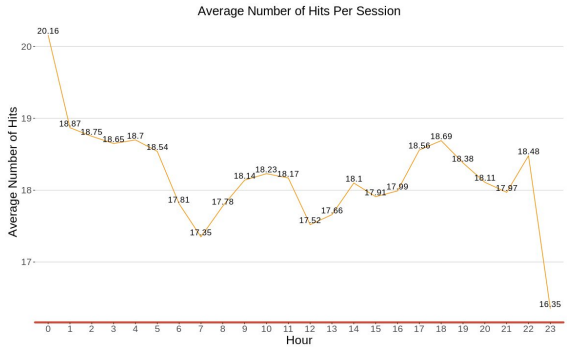

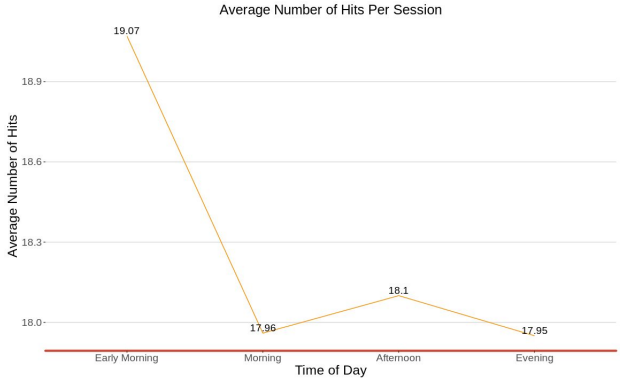

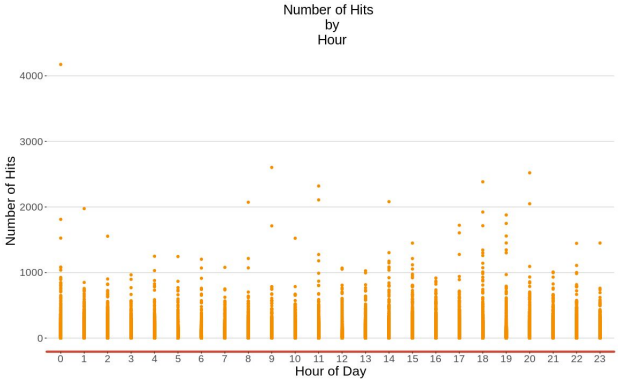

Time of day: Cont'd

- Generally, there is higher traffic on the Trivago Website early in the morning (i.e between midnight and 5:00 am), as compared to other times.
- Traffic is higher during weekend afternoons as compared to weekday afternoons. This may be due to the fact that during weekdays, people are working during the day,that is why they visit the website either during early mornings, or in the evenings(i.e between 7:00 pm and 11:00pm).

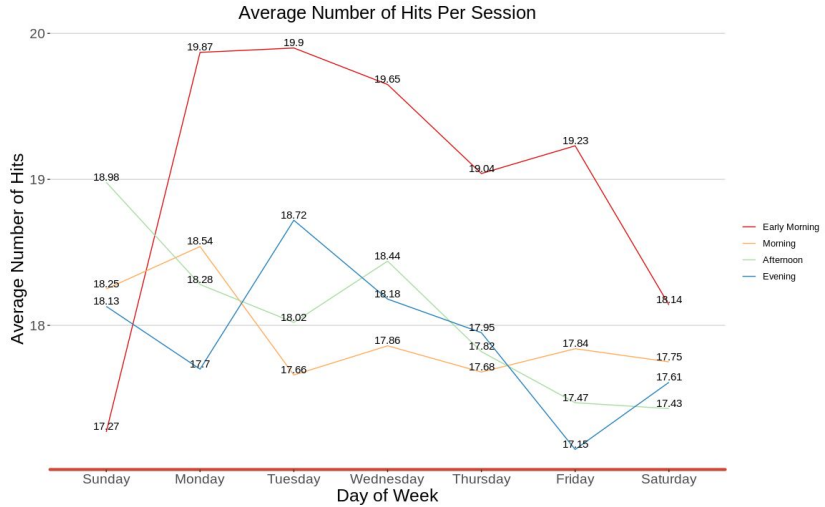

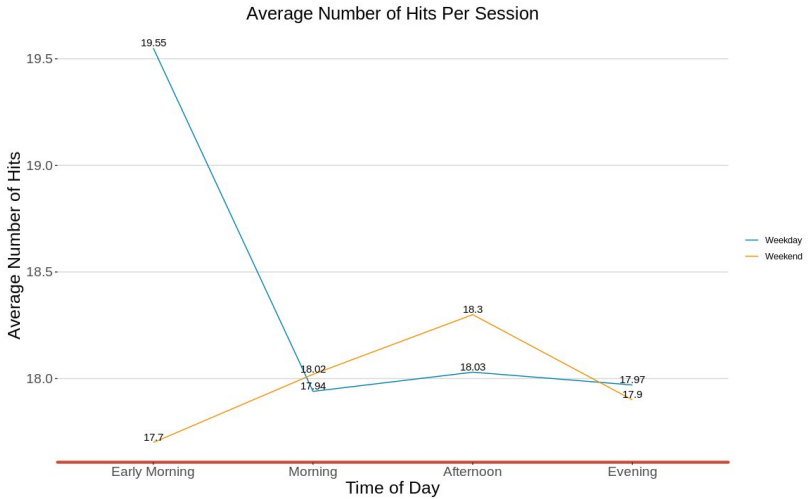

Locations visited

- There is a **moderate positive correlation(0.403462)** between the average number of hits and the number of pages visited per session .
- There is a higher number of hits for the sessions that start with path id 0, as compared to those that start with other path ids.

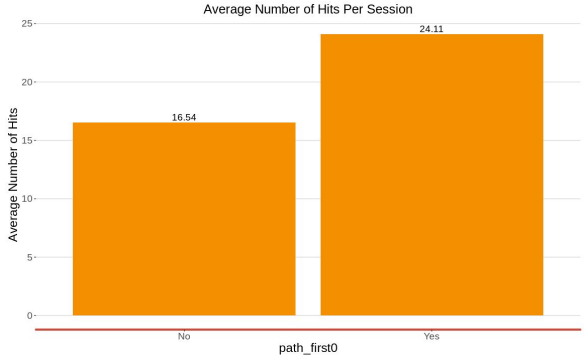

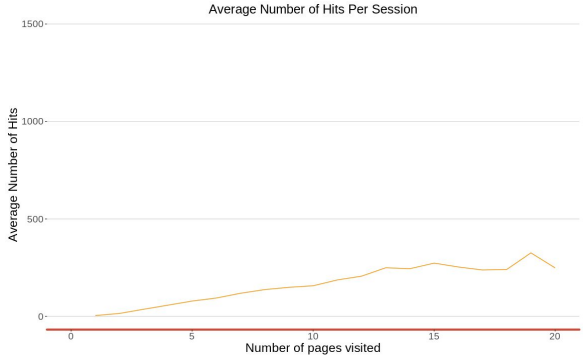


```
             Df    Sum Sq Mean Sq F value             Pr(>F)
path_first0   1   6104841 6104841    4019 <0.0000000000000002 ***
Residuals    619233 940571736    1519
```

```
        Pearson's product-moment correlation

data:  hits_df2$num_locs and hits_df2$hits
t = 346.98, df = 619233, p-value < 0.00000000000000022
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.4013746 0.4055452
sample estimates:
     cor
0.403462
```

Platforms

- L1, L2 and L5 result into a higher number of hits, as opposed to the rest of the platforms.

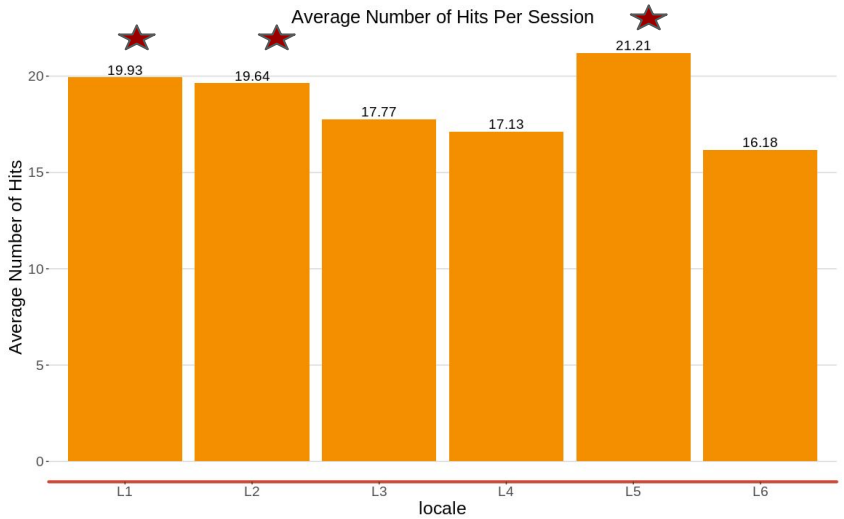

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
locale	5	1545266	309053	202.5	<0.0000000000000002 ***
Residuals	619229	945131311	1526		

Agent IDs

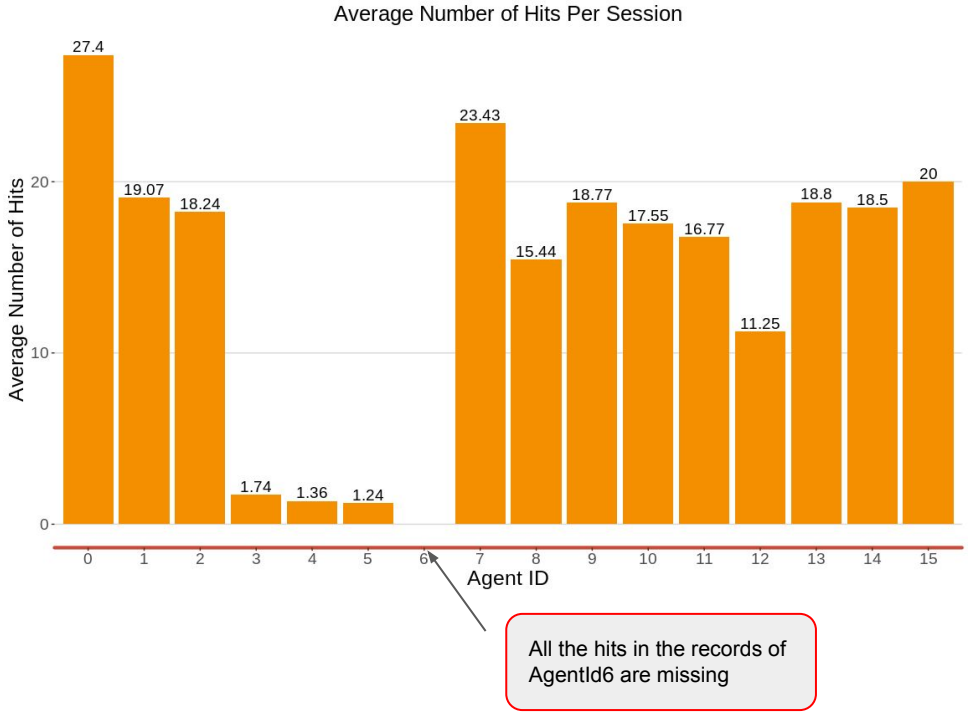

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
agent_id	14	1703390	121671	79.73	<0.0000000000000002 ***
Residuals	619220	944973187	1526		

Entry Page

- The entry pages that result into a higher number of hits are 2202 and 2706.

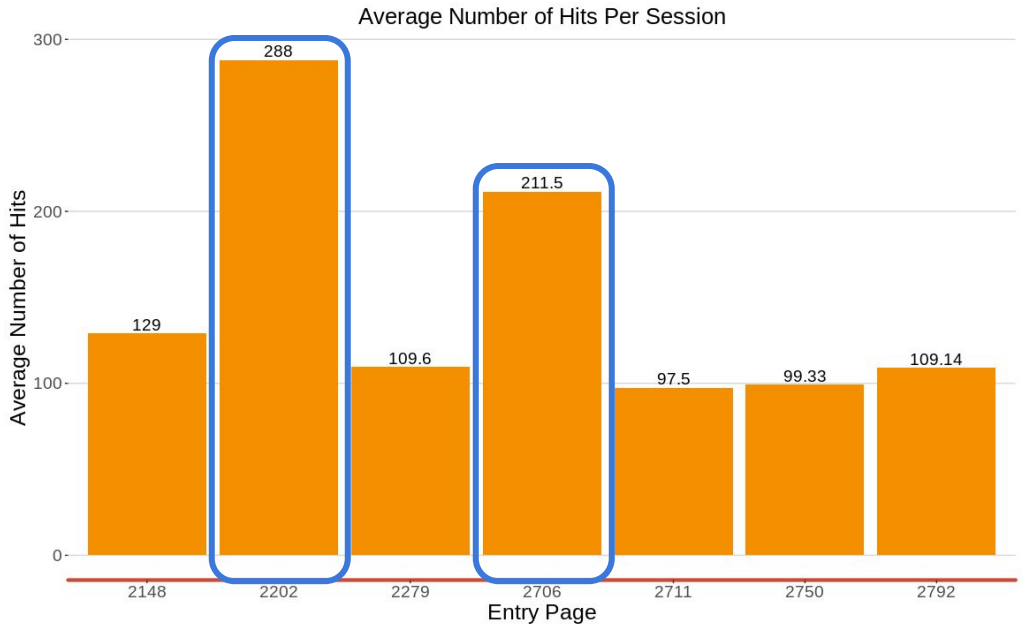

Traffic

- The lesser the value of the traffic type, the higher the number of hits per session.
- I can confidently say that the higher the value of traffic type, the lesser the engagement on the website.

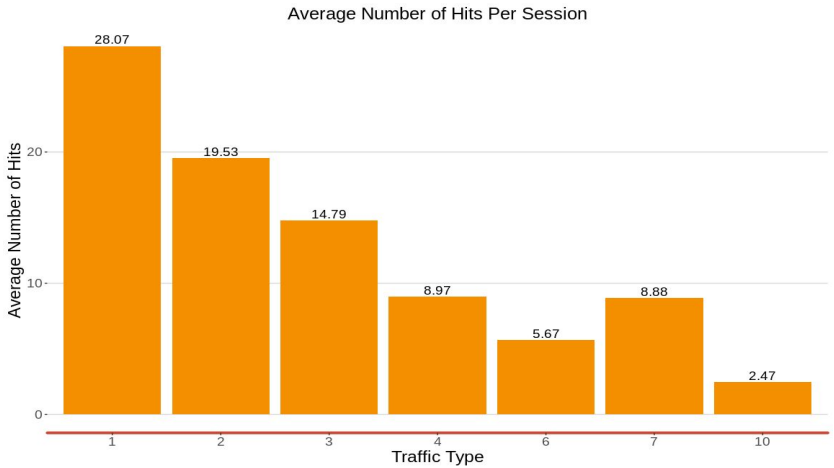

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
traffic_type	6	43481933	7246989	4969	<0.0000000000000002 ***
Residuals	619228	903194644	1459		

Here, I assumed that the traffic type is coded in such a way that 1 represents higher traffic on the website, and 10 represents lesser traffic, as this behavior is depicted on the graph.

Session Duration

- There is a **moderate positive relationship(0.2455381)** between the session duration and the number of hits per session.

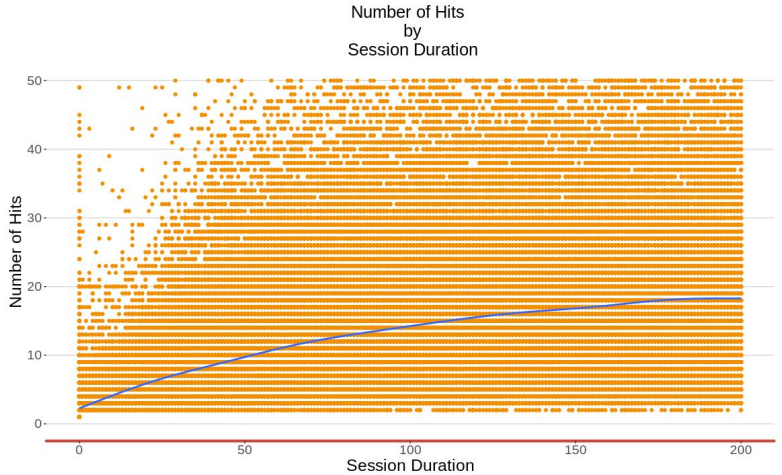


```
        Pearson's product-moment correlation

data:  hits_df2$session_duration and hits_df2$hits
t = 199.25, df = 618819, p-value < 0.00000000000000022
alternative hypothesis: true correlation is not equal to 0
95 percent confidence interval:
 0.2431873 0.2478700
sample estimates:
      cor
0.2455301
```

Predictive Model

Examining Outliers

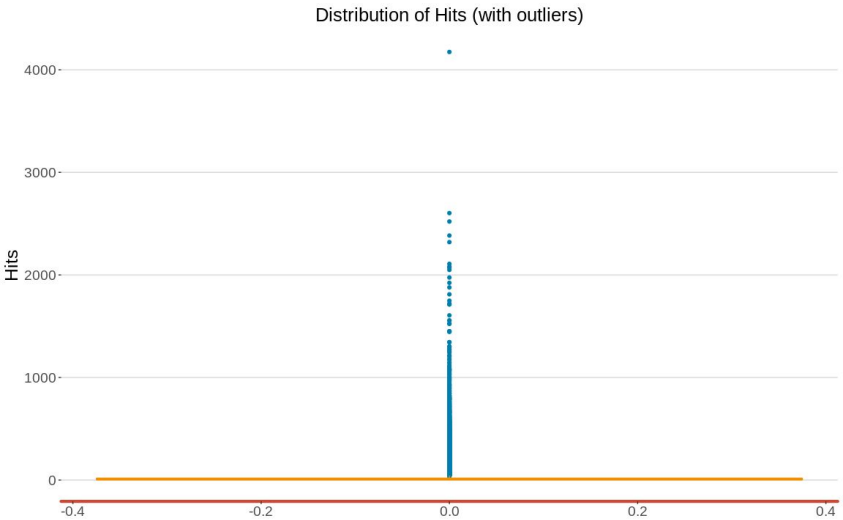

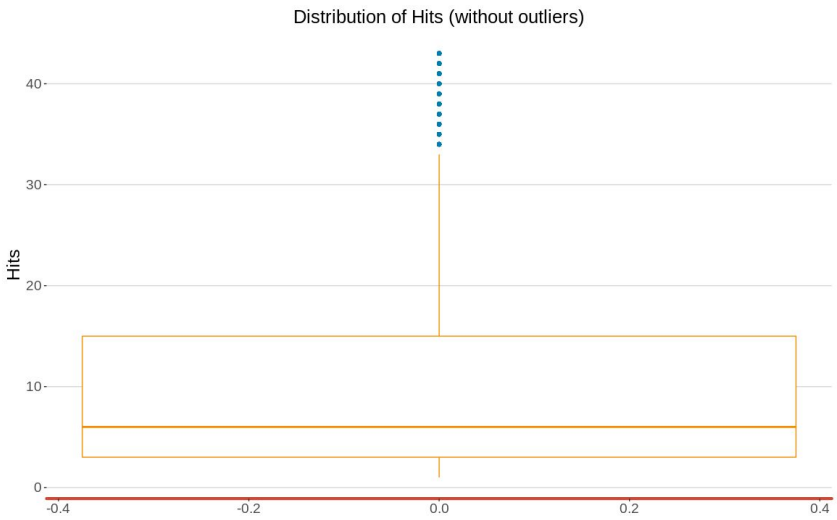

374,732 outliers were dropped at this point

Splitting the data into train and test

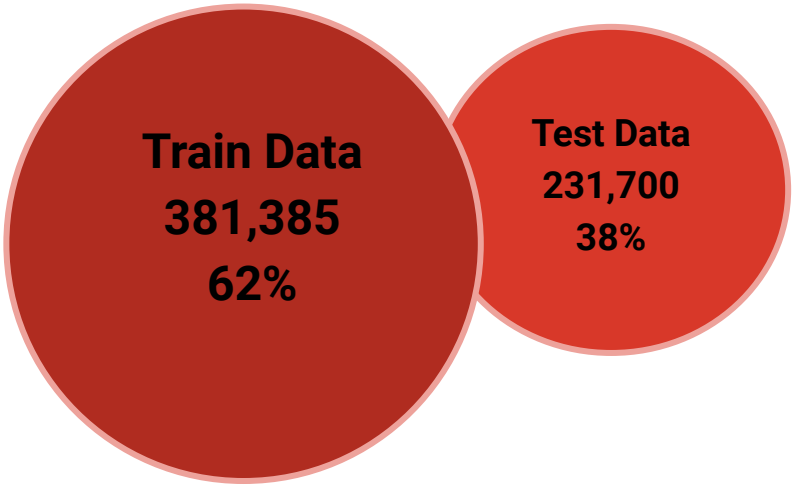

Alert!!

- My laptop’s computation power could not handle the whole dataset, so I reduced the size of my train and test dataset by **5**.

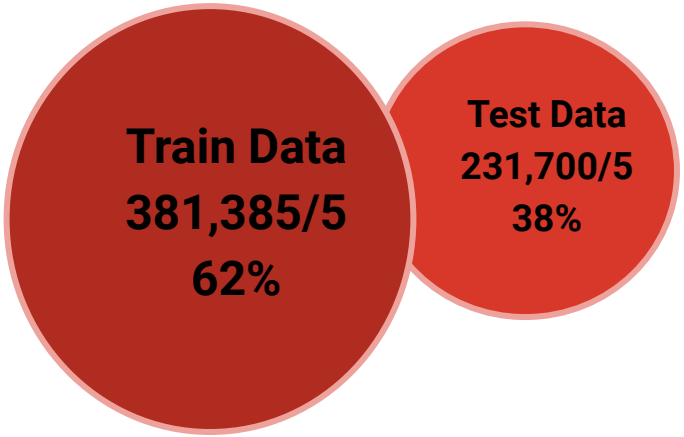

Variable Importance

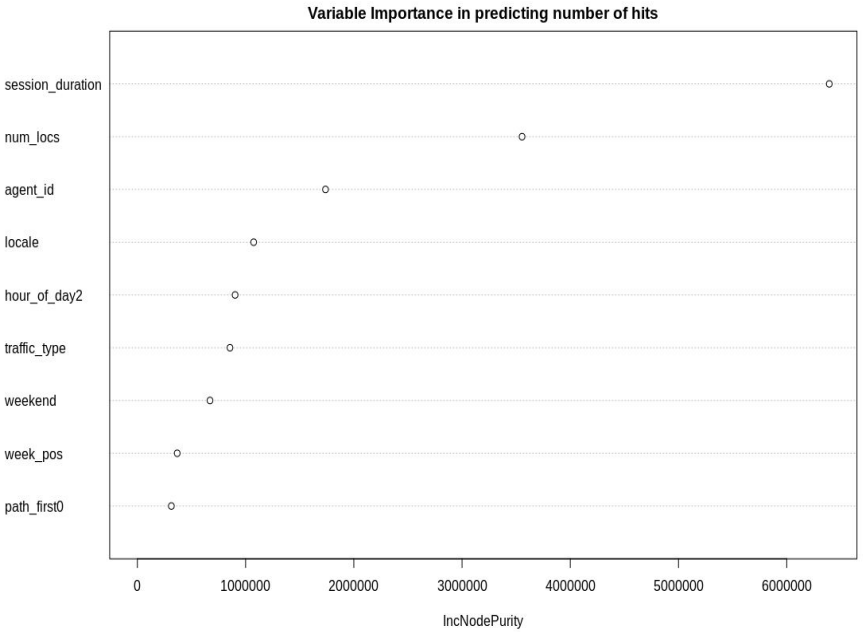

Variables in order of Importance	
1	session_duration
2	num_locs
3	agent_id
4	locale
5	hour_of_day2
6	traffic_type
7	weekend
8	week_pos
9	path_first0

Random Forest

- See data attached together with this presentation deck.

Support Vector Machines

- See data attached together with this presentation deck.

The End!